

COGS 300



Distribution 04

Nov 20/25

①


Last two classes, we've discussed reinforcement learning. Now, we will ask: how does a multi-agent system or network learn?

One evaluative "function" is death. Evolutionary algorithms leverage this concept: each "agent" has a set of parameters that randomly vary, weighted by success.




Cutting edge AI agents basically find a way to force the eval to be in text.

An example might be learning to take the right speed at a turn:



The question is how to apply a reward fn.


Let's sketch a reasonable approach:



- ① How do you detect running into a wall?
- ② How do you punish it?
- ③ How do you reward not running into a wall?

Ex. if $\underline{OL} > 100$ and $\underline{sp_L} < 10 \rightarrow -100$

- ④ How do you actually do this iRL?



ex.

But how do you evaluate the evaluator?

back propagation

let's go back to the 1950s.

The history of CS really is the history of trying to make thinking machines.

(3)


The architecture of modern ML was designed in the 1950s. Backprop. in 1970s.

McCulloch + Pitts developed a calculus of a neuron in 1943.

Frank Rosenblatt built a machine to actually calculate in 1957.

This is the perceptron.

Idea: using circuits, implement the function of a neuron.



This should look like an extension of a logic gate.

$$a \ b \ c \quad \text{out} = a+b+c \geq \tau ?$$

0	0	0	0
0	0	1	0
0	1	0	0
0	1	1	1
1	0	0	0
1	0	1	1
1	1	0	1
1	1	1	1

$$p_{00} \ p_{01} \ p_{02} \dots p_{1k} \quad \text{out} = \sum p_{ij} \geq \tau ?$$



$$\tau = 12$$

both A?



vs



activation

weighted mask



④

* How do you design the weights so that



correct ✓

} only two correct.



correct ✓



incorrect ✗

etc.

any other example incorrect.

Backprop: How do this automatically?

$$\text{scale.}$$

$$\begin{pmatrix} \Sigma & \\ & \end{pmatrix} \Rightarrow z \leq \tau$$

but we know should be 1.

try again! New weights:

$$\begin{pmatrix} \Sigma & \\ & \end{pmatrix} \Rightarrow -2 \leq \tau$$

hmm... need to do better than random.

* How can we evaluate our weights?

loss: how far are we from the goal?

(5)

computing power simply wasn't enough until
~ the 2010s to make this

Deep Blue GOFAT IBM.

Watson NLP ML

unwarrantly, but approaches mix. GOFAT + mL.

* Design an architecture for your
robot to leverage AI.

→ which parts can be mL?

→ which parts can be GOFAT?

→ which parts need UCM eval?

↳ how to construct it?

Warm up: Draw isomorphisms of a random walk.

Eg.



random walk



etc.

random walk's angle



random walk in pixel space.



Tonight - sun: East Van Culture Crawl
→ Van Hack Space on Venables.




(2)

Non-mathy reward functions?

Optical character recognition (OCR)

Img \rightarrow text





state

UL	Ur	UR	EL	ER	ML	MR
Dist L	pist m	Dist R	Speed L	Speed R	out L	out R

action

$$0-25\%$$

$$0-100\%$$

- ① Detect running into a wall
- ② Reward / punish?
- ③ Can you train the robot autonomously?

③

④

① No change to dist.
if encoder < t → not moving

②

int total_dist




loop:

$t_d += \Delta t$ +

reward(total_dist)

③



McCulloch + Pitts.

Frank Rosenblatt 1957


perception

=

(5)



$$\begin{array}{l} \text{a } 1, 0 \\ \text{b } 0, 0 \\ \text{c } 0, 1 \\ \vdots \\ \text{out} = a + b + c > \tau \end{array}$$




$$p_{00} \quad p_{10} \quad p_{20} \dots \quad p_{01} \dots p_{ij}$$

1 1 1 0

$$\text{out} = \sum p_{ij} > \tau$$

$$\tau = 12$$

weighted mask




⑥



come up with a mask
that correctly classify

-1, 1



err?
=

loss